

RAG-SERVER Test Document

Page 1 of 3

This is a sample PDF document for testing the PDF loader.
It contains multiple pages with different content.

Milvus is an open-source vector database built for AI applications.
It supports billion-scale vector search with high performance.

Dense vectors represent semantic meaning of text.

Sparse vectors like BM25 are good for keyword matching.

Hybrid search combines both for better results.

RRF (Reciprocal Rank Fusion) is an algorithm for merging results.

Page 3 of 3

Document ingestion pipeline includes:

1. PDF loading
2. Text splitting
3. Embedding generation
4. Vector storage

End of test document.